

CMSC4980: Final Report
Predicting the Directional Movement of Stock Prices
Albert Koy & Nadeem Malik

Project repository: https://github.com/akoy93/CMSC4980_Final_Project

Collected/transformed data:

https://github.com/akoy93/CMSC4980_Final_Project/tree/master/data

Model Predictions:

https://github.com/akoy93/CMSC4980_Final_Project/tree/master/data/2014-12-02_1yr_predictions

Introduction

For our final project, we created a model to predict the directional movement of a stock price given historical data, what is known as technical analysis. The goal of our project is to predict whether a stock will close higher than it opened given 10 prior days of stock data. We applied machine learning algorithms to historical stock data and we found that by utilizing logistic regression, we were able to predict correctly with ~81% accuracy and ~86% AUC with cross-validation. Our intermediate report focused on the research and survey of commonly used machine learning techniques for technical analysis, which led us to use logistic regression. This report will focus on the practical aspect of our project, the process for getting the data, creating the models, and displaying our results.

Process and Model Development

Gathering and Preprocessing Data

The first challenge in taking on this project was to find a source for historical stock market data. Unfortunately, we couldn't find many large, free, downloadable sets of historical stock market data on the Internet, so we had to write a web scraper to gather this data. We did this by writing a Ruby script to scrape the Yahoo! Finance page to grab data for all U.S. stocks on a daily timeframe over the past 5 years. We now have cleanly formatted CSV files for each U.S. Stock symbol (*Figure 2*) containing the Open, High, Low, Close, and Adjusted Close prices for each stock along with the daily volume (*Figure 3*).





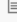


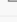
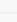
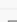
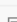





| ⚠ Sorry, we had to truncate this directory to 1,000 files. 4,878 entries were omitted from the list. | | |
|--|---------------------------------------|-------------|
|  A.csv | finished script to compute technicals | an hour ago |
|  AA.csv | finished script to compute technicals | an hour ago |
|  AAC.csv | finished script to compute technicals | an hour ago |
|  AAL.csv | finished script to compute technicals | an hour ago |
|  AAMC.csv | finished script to compute technicals | an hour ago |
|  AAME.csv | finished script to compute technicals | an hour ago |
|  AAN.csv | finished script to compute technicals | an hour ago |
|  AAOI.csv | finished script to compute technicals | an hour ago |
|  AAON.csv | finished script to compute technicals | an hour ago |
|  AAP.csv | finished script to compute technicals | an hour ago |
|  AAPL.csv | finished script to compute technicals | an hour ago |
|  AAT.csv | finished script to compute technicals | an hour ago |
|  AAU.csv | finished script to compute technicals | an hour ago |
|  AAV.csv | finished script to compute technicals | an hour ago |
|  AAVL.csv | finished script to compute technicals | an hour ago |
|  AAWW.csv | finished script to compute technicals | an hour ago |

Figure 2

| Date | Open | High | Low | Close | Volume | Adj Close |
|------------|-------|-------|-------|-------|---------|-----------|
| 2014-11-18 | 40.25 | 41.16 | 40.05 | 40.80 | 5891400 | 40.80 |
| 2014-11-17 | 41.57 | 41.70 | 41.21 | 41.24 | 2696100 | 41.24 |
| 2014-11-14 | 41.36 | 41.65 | 41.29 | 41.57 | 1331300 | 41.57 |
| 2014-11-13 | 41.55 | 41.74 | 41.26 | 41.45 | 1805500 | 41.45 |
| 2014-11-12 | 41.30 | 41.54 | 41.02 | 41.45 | 2504300 | 41.45 |
| 2014-11-11 | 41.52 | 42.11 | 41.34 | 41.66 | 2163800 | 41.66 |
| 2014-11-10 | 40.98 | 41.55 | 40.88 | 41.53 | 2583900 | 41.53 |
| 2014-11-07 | 41.39 | 41.48 | 40.86 | 40.93 | 2458700 | 40.93 |
| 2014-11-06 | 40.39 | 41.41 | 40.26 | 41.37 | 2968300 | 41.37 |
| 2014-11-05 | 40.61 | 40.67 | 39.98 | 40.13 | 2341400 | 40.13 |
| 2014-11-04 | 40.26 | 40.83 | 39.84 | 40.18 | 5765200 | 40.18 |
| 2014-11-03 | 39.90 | 41.24 | 39.49 | 40.84 | 4559800 | 40.84 |

Figure 3

After retrieving the raw stock data, the next step was to compute technical analysis indicators for each stock. This meant writing a script to transform each of the 5000+ CSV files into a file that would give us some more information to work with for each stock. We used Python to write this script and relied heavily on the NumPy, Pandas, and TA-Lib libraries. TA-Lib is an

industry standard library for performing technical analysis, with many functions to help compute indicators. To augment the raw data that we retrieved from Yahoo! Finance, we computed the 50 Day Average Volume, 30 Day, 4 Week, 10 Week, and 30 Week Simple Moving Averages, the Daily 20.2 Upper, Middle, and Lower Bollinger Bands, the Daily 12.26.9 MACD, 12.26.9 MACD Signal, and 12.26.9 MACD Histogram values, the Relative Strength Index with 14 bars as the timeframe, and the Daily 10.4 and 10.4.4 Stochastic values.

| | Date | Open | High | Low | Close | Volume | Adj Close | 50DayAvgVol | 30DaySMA |
|----|------------|-------|-------|-------|-------|---------|-----------|-------------|---------------|
| 0 | 2014-11-18 | 40.25 | 41.16 | 40.05 | 40.8 | 5891400 | 40.8 | 3203320.0 | 48.618 |
| 1 | 2014-11-17 | 41.57 | 41.7 | 41.21 | 41.24 | 2696100 | 41.24 | 3121452.0 | 49.0913333333 |
| 2 | 2014-11-14 | 41.36 | 41.65 | 41.29 | 41.57 | 1331300 | 41.57 | 3105158.0 | 49.6023333333 |
| 3 | 2014-11-13 | 41.55 | 41.74 | 41.26 | 41.45 | 1805500 | 41.45 | 3103900.0 | 50.1136666667 |
| 4 | 2014-11-12 | 41.3 | 41.54 | 41.02 | 41.45 | 2504300 | 41.45 | 3089896.0 | 50.593 |
| 5 | 2014-11-11 | 41.52 | 42.11 | 41.34 | 41.66 | 2163800 | 41.66 | 3081134.0 | 51.0846666667 |
| 6 | 2014-11-10 | 40.98 | 41.55 | 40.88 | 41.53 | 2583900 | 41.53 | 3094214.0 | 51.5953333333 |
| 7 | 2014-11-07 | 41.39 | 41.48 | 40.86 | 40.93 | 2458700 | 40.93 | 3097894.0 | 52.116 |
| 8 | 2014-11-06 | 40.39 | 41.41 | 40.26 | 41.37 | 2968300 | 41.37 | 3087736.0 | 52.635 |
| 9 | 2014-11-05 | 40.61 | 40.67 | 39.98 | 40.13 | 2341400 | 40.13 | 3061726.0 | 53.1523333333 |
| 10 | 2014-11-04 | 40.26 | 40.83 | 39.84 | 40.18 | 5765200 | 40.18 | 3039650.0 | 53.7466666667 |

| 4WeekSMA | 10WeekSMA | 30WeekSMA | 20.2UpperBB | 20.2MiddleBB |
|----------|-----------|---------------|---------------|--------------|
| 46.3355 | 52.0862 | 55.1544666667 | 59.2177753813 | 46.3355 |
| 46.9955 | 52.4192 | 55.2403333333 | 60.0274292125 | 46.9955 |
| 47.551 | 52.7474 | 55.3194666667 | 60.5011288024 | 47.551 |
| 48.086 | 53.0744 | 55.3933333333 | 60.8867912255 | 48.086 |
| 48.599 | 53.3992 | 55.4688 | 61.1140691568 | 48.599 |
| 49.116 | 53.7314 | 55.5528 | 61.2557139999 | 49.116 |
| 49.6225 | 54.0596 | 55.6455333333 | 61.3125878953 | 49.6225 |
| 50.1855 | 54.3722 | 55.7330666667 | 61.3344541662 | 50.1855 |
| 50.821 | 54.6966 | 55.8225333333 | 61.2103019977 | 50.821 |
| 51.504 | 55.0168 | 55.9172 | 61.0826479213 | 51.504 |

| 20.2LowerBB | 12.26.9MACD | 12.26.9MACDSignal | 12.26.9MACDHist | RSI14 | 10.4STOC | 10.4.4STOC |
|---------------|----------------|-------------------|-----------------|---------------|---------------|---------------|
| 33.4532246187 | -3.54979768637 | -3.3683180917 | -0.181479594666 | 24.9545947239 | 47.9681330397 | 32.8918681589 |
| 33.9635707875 | -3.61532868011 | -3.32294819304 | -0.292380487076 | 25.8518585235 | 41.4216093414 | 23.88742912 |
| 34.6008711976 | -3.6911834586 | -3.24985307127 | -0.441330387331 | 26.5158331494 | 29.4107704553 | 16.4882139706 |
| 35.2852087745 | -3.76791179829 | -3.13952047444 | -0.628391323851 | 25.8729693976 | 12.766959799 | 11.491048995 |
| 36.0839308432 | -3.79882088933 | -2.98242264347 | -0.816398245859 | 25.8729693976 | 11.9503768844 | 10.1248429648 |
| 36.9762860001 | -3.7830165526 | -2.77832308201 | -1.00469347059 | 26.2190770856 | 11.8247487437 | 8.92744974874 |
| 37.9324121047 | -3.72920059478 | -2.52714971436 | -1.20205088042 | 25.6473340401 | 9.42211055276 | 13.0948268126 |
| 39.0365458338 | -3.59303258177 | -2.22663699426 | -1.36639558751 | 23.0931934421 | 7.30213567839 | 23.2726019073 |
| 40.4316980023 | -3.30587948035 | -1.88503809738 | -1.42084138297 | 23.6463252806 | 7.1608040201 | 39.6870879476 |
| 41.9253520787 | -2.9360821839 | -1.52982775164 | -1.40625443227 | 18.5404400456 | 28.4942569993 | 61.8440795085 |

Figure 4

Developing Models

Next, once we had our raw stock data, we set out to train three types of models that we identified in our intermediate report: Logistic Regression, SVM with Linear kernel, and Random Forest Classifier. We wanted to be able to input 10 days of data into our model in order to predict the directional movement of the last day's closing price. The biggest challenge in this step was feature selection. Initially, we found that by training on all of the 21 features, we were getting very poor results. In most cases, we were getting predictive accuracy that was essentially the same as flipping a coin for this binary classification problem. Moreover, since we had 5 years of data for ~6000 stocks and 21 features per day, computation time was extremely expensive and we were never able to train our models with a non-trivial set of our raw data.

One major insight that we made that allowed us to speed up our training process was that the technical indicators we calculated were all derivatives of the initial set of raw data that we started with. Thus, we decided to start training on our initial set of data without the technical indicators, which allowed us to develop a model, but still without any luck in increasing predictive accuracy.

The next major insight we made was that was that we had to normalize our data to add consistency to our dataset between different stocks. We decided to normalize the Open, High, Low, and Close prices by dividing by the value of the Open price on the first day in the 10-day window. Additionally, we decided to normalize the Volume by dividing by the 50-day average volume. At this point, we were able to finish training our Logistic Regression model with a large number of stocks, but the SVM and Random Forest Classifiers were still too slow. We were able to achieve decent accuracy with our Logistic Regression model, predicting ~70% of outcomes correctly with a random sample of data which was roughly half of our dataset.

In order to get a more representative dataset, we observed that it would be worthwhile to train our models by using a market index, such as the S&P 500. The S&P 500 is commonly used to evaluate the general stock market trend, as it is believed to be one of the most representative

stock indexes of the entire stock market. By using the S&P 500 as our training set, we were able to achieve a small boost in predictive accuracy, and since this was less data than we were using originally, we were able to finish training our SVM and Random Forest models.

The last insight we made was inspired by a talk given in the University of Maryland's Stock Market Class (HONR348M). According to one guest speaker in that class who makes his living trading stocks in the stock market, he uses only the price action of a stock and draws trend lines between price bars in order to determine whether or not to enter a stock. He mentioned that he didn't use trading volume as an indicator, which is very important according to most expert traders. With that in mind, we decided to train our model without using normalized volume as a feature, and we achieved a ~86% AUC on 5-fold cross-validation when training on 5 years of the S&P 500.

Finally, with more testing, we found that including the High and Low prices in our training data did not make much of a difference in predictive accuracy, so we excluded them from our training set in order to speed up the training process. We ended up with a Logistic Regression model that takes in 19 points of data, the Open and Close for the previous 10 days except for the Close on the last day, and predicts whether the next Close will be higher than the observed Open. We observed that the Linear SVM was slower to train as well as slightly less accurate, and the Random Forest Classifier was nowhere near as accurate, so we decided to use the Logistic Regression method for our final model.

Predictions

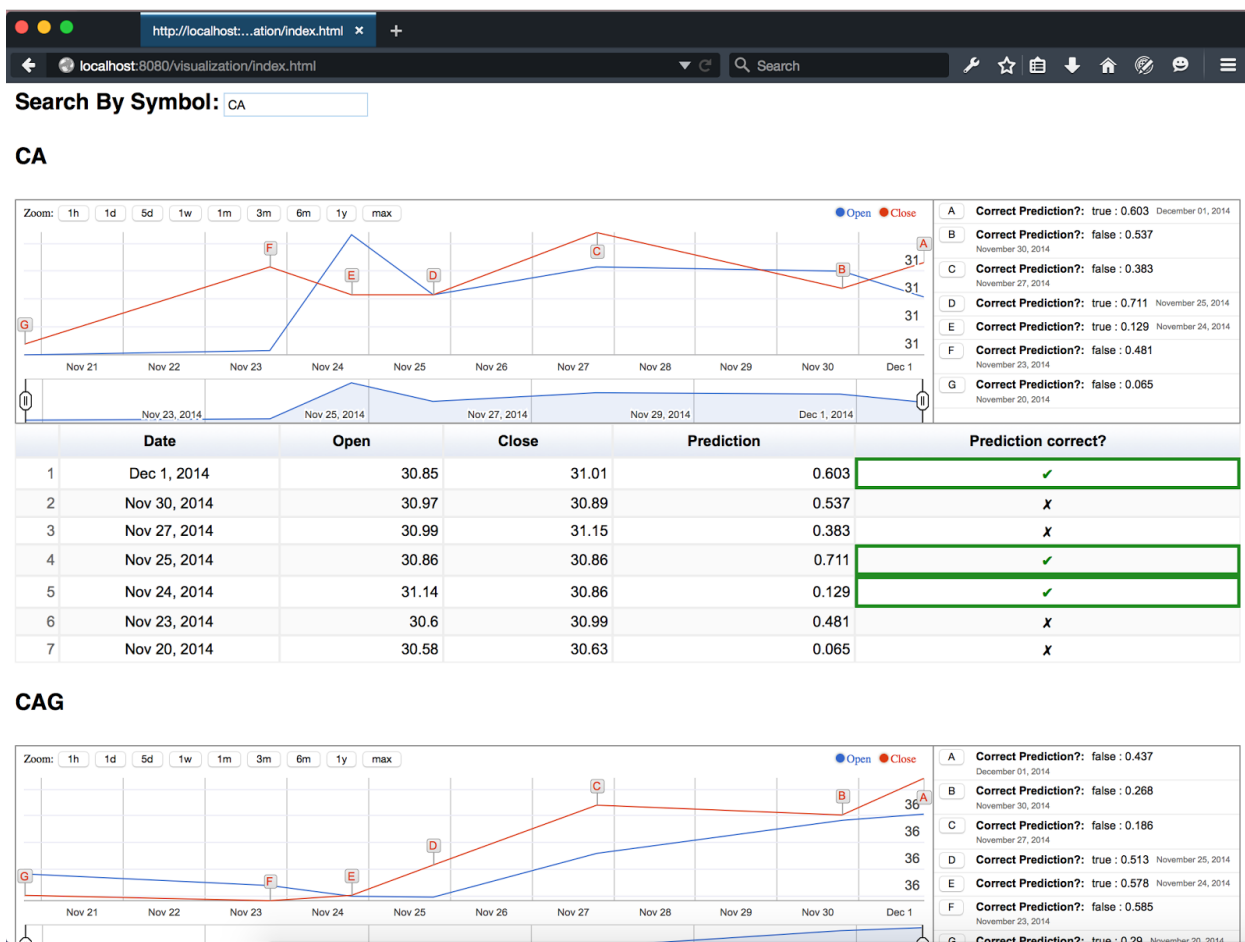
A set of all prediction files is available here:

https://github.com/akoy93/CMSC4980_Final_Project/tree/master/data/2014-12-02_1yr_predictions

We were able to achieve 91% AUC with cross validation when training our Logistic Regression model with S&P 500 stock data over 5 years. This translated to ~81% correctness when making predictions across all ~6000 U.S. stocks. Below is a sample prediction file that is output by our model. It includes "Prediction" and "Prediction Correct" columns appended to the raw stock data which, respectively, give a probability from 0 to 1 of whether the stock will Close higher than the open and a boolean value indicating whether or not the prediction was correct. For our purposes, a value over 0.5 indicates a stock will Close higher, a value below 0.5 indicates a stock will close lower.

| Date | Open | High | Low | Close | Volume | Adj Close | Prediction | Prediction Correct |
|------------|------|------|------|-------|--------|-----------|-------------------|--------------------|
| 2014-11-18 | 1.15 | 1.18 | 1.1 | 1.15 | 234400 | 1.15 | 0.0438078903889 | False |
| 2014-11-17 | 1.1 | 1.18 | 1.08 | 1.14 | 130000 | 1.14 | 0.995677314033 | True |
| 2014-11-14 | 1.05 | 1.13 | 1.03 | 1.13 | 458700 | 1.13 | 0.998264290456 | True |
| 2014-11-13 | 1.13 | 1.13 | 1.04 | 1.06 | 87200 | 1.06 | 0.000472620449082 | True |
| 2014-11-12 | 1.13 | 1.13 | 1.08 | 1.13 | 84800 | 1.13 | 0.246890176774 | False |
| 2014-11-11 | 1.08 | 1.12 | 1.05 | 1.12 | 102900 | 1.12 | 0.995661375372 | True |
| 2014-11-10 | 1.05 | 1.07 | 1.02 | 1.07 | 108600 | 1.07 | 0.980921540854 | True |
| 2014-11-07 | 1.04 | 1.08 | 1.0 | 1.05 | 238400 | 1.05 | 0.797113489196 | True |
| 2014-11-06 | 1.0 | 1.04 | 0.97 | 1.03 | 237600 | 1.03 | 0.987279742647 | True |
| 2014-11-05 | 1.09 | 1.09 | 0.99 | 1.04 | 392400 | 1.04 | 0.000192821823056 | True |

Visualization



(Screenshot shows old prediction data. Newer data is more accurate.)

We used d3.js and Google Visualization API to create a visualization demo that shows the results of our application. The demo displays a series of stocks with corresponding tables and charts. The table lists the open and close prices by date, our prediction estimate, and whether our prediction was correct for that day or not. The chart for each stock shows annotations for when our predictions were correct and has controls to adjust the graphic display by range of time selected. To come up with this visualization, we first read stock symbols from a text file. For each stock symbol, we read its corresponding CSV file to get a certain number of entries to display. Next, we created a table and chart element in our web page, for each set of entries. There were several challenges we faced when writing code for visualizations. For example, when we tried to use the d3.js library to parse in CSV file for each stock symbol, it was impossible to load all the files as any browser we tried would crash due to insufficient resources. We learned that data visualization has significant resources challenges that need to be realized before one comes up with a demo. On another note, we found the Google Visualization API to be very simple to use and quite elegant. This API uses a DataTable object to store data that is read in from JavaScript arrays, which is then used to create the chart or table you would like to display.

Our Experiences

When initially starting this project, we believed that most of our efforts would be spent choosing different machine learning algorithms and configurations as well as computing our own custom indicators to increase our predictive accuracy. This turned out to not be the case at all. In reality, the vast majority of our time was spent gathering data, writing scripts to process the data, and figuring out how to reduce the amount of data that we needed in order to train the models we used for our project. Given the limited hardware we have, we were required to abandon the vast majority of our data when training our models, which was especially challenging to decide. The actual machine learning implementations required ~10 lines of code by using the SK-Learn Python library, which was very surprising to us. In addition, we realized that visualizations can be more complicated than they seem, even with the vast amount of APIs that make it seem easy, there are other challenges that need to be considered, such as graphical rendering and performance. If we had more time, we would be able to further optimize our visualization for better performance. Overall, this project was a great experience in exploring the realm of practical machine learning, and we hope to have further opportunities involving such work in the future.

Code:

Scraper - get_stocks.rb (Scraping raw data from Yahoo! Finance)

```
require 'date'
require 'open-uri'
```

```

# usage ruby get_stocks.rb #{STOCK_SYMBOLS} #{NUMBER_OF_YEARS}

unless ARGV.length == 2
  puts "Incorrect Usage. Use: ruby get_stocks.rb {STOCK_SYMBOLS_FILE} {NUMBER_OF_YEARS}"
else
  current_date = Date.parse(Time.now.to_s)
  filename = ARGV[0]
  num_years = ARGV[1].to_i
  directory = "#{current_date.to_s}_#{num_years}yr"
  symbols = []

  # remove existing directory
  system("rm -rf #{directory}")
  system("mkdir #{directory}")

  # read stock symbols from file
  File.foreach(filename) do |sym|
    sym.chomp!
    puts "Fetching #{sym}..."
    # api url for stock data retrieval
    retrieval_url = "http://real-chart.finance.yahoo.com/table.csv?s=#{sym}" +
      "&a=#{current_date.month - 1}&b=#{current_date.day}&c=#{current_date.year -
num_years}&d=#{current_date.month - 1}" +
      "&e=#{current_date.day}&f=#{current_date.year}&g=d&ignore=.csv"
    begin
      open("#{directory}/#{sym}.csv", "wb") do |file|
        file << open(retrieval_url).read
      end
      symbols << sym
    rescue
      puts "Unable to retrieve #{sym}"
    end
  end
end

```



```

File.open("#{directory}/symbol_list.txt", "wb") do |file|
  symbols.each { |sym| file.puts sym }
end

puts "DONE"
end

```

Prediction - predict.py (Using our model to output predictions)

```

# Usage: python generate_training.py {DIRECTORY} {MODEL_FILE_NAME}
import numpy as np
import pandas as pd
import glob
import sys
import pickle
from subprocess import call

NUM_DAYS_IN_WINDOW = 10
NUM_POINTS_PER_DAY = 2

if len(sys.argv) != 3:
    print "Incorrect usage. Use: \"python generate_training.py {DIRECTORY} {MODEL_FILE_NAME}\""
    sys.exit()

directory = sys.argv[1]
new_directory = "%s_predictions" % directory
model_file_name = sys.argv[2]

```

```

# create new directory
call(["rm", "-rf", new_directory])
call(["mkdir", new_directory])

with open(model_file_name, "rb") as f:
    model = pickle.load(f)
    print "Model loaded."

# get all csv files in the given directory
files = glob.glob("%s/*.csv" % directory)
file_num = 0

for f in files:
    file_num += 1
    print("(%d/%d) - Making predictions for %s..." % (file_num, len(files), f[len(directory) + 1:])
    try:
        data = pd.read_csv(f)[::-1]
        data = data.set_index('Date')
        data['Prediction'] = np.nan
        data['Prediction Correct'] = np.nan

# initialize window
window = [data.loc[[0]]['Open']]

i = 0
for i in range(1, NUM_DAYS_IN_WINDOW):
    row = data.loc[[i]]
    window = [row['Open'], row['Close']] + window

```

```

# normalize

window_arr = np.hstack(np.asarray(window))
normalized = window_arr / window_arr[0]

# result

idx = data.loc[[0]].index
data.loc[idx, 'Prediction'] = model.predict_proba([normalized])[0,1]
data.loc[idx, 'Prediction Correct'] = bool(model.predict([normalized])[0]) == (data.loc[idx, 'Close'] >=
data.loc[idx, 'Open'])

while i < data.shape[0] - 1:
    i += 1
    row = data.loc[[i]]
    window_arr = window_arr[:-NUM_POINTS_PER_DAY]
    window_arr = np.hstack([np.hstack([row['Open'], row['Close']]), window_arr])
    normalized = window_arr / window_arr[0]
    idx = data.loc[[i - NUM_DAYS_IN_WINDOW + 1]].index
    data.loc[idx, 'Prediction'] = model.predict_proba([normalized])[0,1]
    data.loc[idx, 'Prediction Correct'] = bool(model.predict([normalized])[0]) == (data.loc[idx, 'Close'] >=
data.loc[idx, 'Open'])

data = data.dropna()[::-1]
data.to_csv(f.replace(directory, new_directory), sep=',', encoding='utf-8')
except:
    print "Unable to make predictions for %s." % f[len(directory) + 1:]

```

Python - generate_training.py (Generating formatted training data to use in our model)

```

# Usage: python generate_training.py {NUMBER_OF_DAYS} {NUM_DAYS_IN_WINDOW} {DIRECTORY}
{OUTPUT_FILE}

import numpy as np
import pandas as pd

```

```

import glob
import sys

if len(sys.argv) != 5:
    print "Incorrect usage. Use: \"python generate_training.py {NUMBER_OF_DAYS} {NUM_DAYS_IN_WINDOW} {DIRECTORY} {OUTPUT_FILE}\""
    sys.exit()

num_days_to_train = int(sys.argv[1])
num_days_in_window = int(sys.argv[2])
directory = sys.argv[3]
output_file = sys.argv[4]

# get all csv files in the given directory
files = glob.glob("%s/*.csv" % directory)
i = 0

results = []

# process all files
for f in files:
    i += 1
    print "(%d/%d) - Generating training data from %s..." % (i, len(files), f[len(directory) + 1:])
    try:
        data = pd.read_csv(f)

        if data.shape[0] >= num_days_to_train:
            # drop unnecessary columns and move into one line
            data = data.set_index('Date')
            data = data.drop('Adj Close', 1)

```

```

data = data.drop('High', 1)
data = data.drop('Low', 1)
data = data.drop('Volume', 1)

# convert into numpy array
data = np.array(data)
data = data[0:num_days_to_train]

result = []
for r in range(num_days_in_window):
    result.append(data[r:(r + data.shape[0] - num_days_in_window + 1),:])

results.append(np.hstack(result))
except:
    print "ERROR: Unable to generate training data from %s" % f[len(directory) + 1:]

results = np.vstack(results)
np.savetxt(output_file, results, delimiter=",")

```

Python - train_models.py (Training a Logistic Regression model)

```

# Usage: python train_models.py {TRAINING_FILE_NAME} {MODEL_OUTPUT_FILE}
# We're training a model to determine if a stock will close above its open on the next day.
# We're given the stock's data for a variable window and its opening price on the last day.

import numpy as np
import sys
import pickle

```

```

import sklearn.linear_model as lm
from sklearn import cross_validation

NUMBER_OF_DAYS_IN_WINDOW = 10
NUM_DATA_POINTS_PER_DAY = 2

def tied_rank(x):
    """
    This function is by Ben Hamner and taken from
    https://github.com/benhamner/Metrics/blob/master/Python/ml_metrics/auc.py
    Computes the tied rank of elements in x.
    This function computes the tied rank of elements in x.
    Parameters
    -----
    x : list of numbers, numpy array
    Returns
    -----
    score : list of numbers
        The tied rank of each element in x
    """
    sorted_x = sorted(zip(x, range(len(x))))
    r = [0 for k in x]
    cur_val = sorted_x[0][0]
    last_rank = 0
    for i in range(len(sorted_x)):
        if cur_val != sorted_x[i][0]:
            cur_val = sorted_x[i][0]
            for j in range(last_rank, i):
                r[sorted_x[j][1]] = float(last_rank+1+i)/2.0
            last_rank = i
        if i==len(sorted_x)-1:
            for j in range(last_rank, i+1):
                r[sorted_x[j][1]] = float(last_rank+i+2)/2.0
    return r

```

```
def auc(actual, posterior):
```

```
    """
```

This function is by Ben Hamner and taken from

https://github.com/benhamner/Metrics/blob/master/Python/ml_metrics/auc.py

Computes the area under the receiver-operator characteristic (AUC)

This function computes the AUC error metric for binary classification.

Parameters

```
-----
```

actual : list of binary numbers, numpy array

The ground truth value

posterior : same type as actual

Defines a ranking on the binary numbers, from most likely to

be positive to least likely to be positive.

Returns

```
-----
```

score : double

The mean squared error between actual and posterior

```
    """
```

```
    r = tied_rank(posterior)
```

```
    num_positive = len([0 for x in actual if x==1])
```

```
    num_negative = len(actual)-num_positive
```

```
    sum_positive = sum([r[i] for i in range(len(r)) if actual[i]==1])
```

```
    auc = ((sum_positive - num_positive*(num_positive+1)/2.0) / (num_negative*num_positive))
```

```
    return auc
```

```
def auc_scorer(estimator, X, y):
```

```
    predicted = estimator.predict_proba(X)[:,1]
```

```
    return auc(y, predicted)
```

```
if len(sys.argv) != 3:
```

```
print "Incorrect usage. Use: \"python train_models.py {TRAINING_FILE_NAME}  
{MODEL_OUTPUT_FILE}\""  
sys.exit()
```

```
num_days_in_window = NUMBER_OF_DAYS_IN_WINDOW  
points_per_day = NUM_DATA_POINTS_PER_DAY  
points_in_window = num_days_in_window * points_per_day
```

```
# load training data
```

```
training = np.genfromtxt(sys.argv[1], delimiter=";")
```

```
# normalize
```

```
for row in training:
```

```
    opening = row[0]
```

```
    for i in range(points_in_window):
```

```
        row[i] = row[i] / opening
```

```
# pull X and y
```

```
X = np.array([p[range(points_in_window - 1)] for p in training])
```

```
y = (training[:, points_in_window - points_per_day + 1] > training[:, points_in_window - points_per_day]) +  
0
```

```
# build model and save
```

```
model = lm.LogisticRegression(penalty = "l2").fit(X, y)
```

```
with open(sys.argv[2], 'wb') as f:
```

```
    pickle.dump(model, f)
```

```
    print("Model saved.")
```



```
# print cross validation score

cv_score = np.mean(cross_validation.cross_val_score(model, X, y, cv=5, scoring = auc_scorer))

print "Cross Validation Score = %f" % cv_score
```

Python - test_model.py (Testing a our models)

```
import numpy as np
import sys
import pickle

NUMBER_OF_DAYS_IN_WINDOW = 10
NUM_DATA_POINTS_PER_DAY = 2

if len(sys.argv) != 3:
    print "Incorrect usage. Use: \"python test_model.py {MODEL_FILE_NAME} {TESTING_FILE_NAME}\""
    sys.exit()

model_file = sys.argv[1]
test_file = sys.argv[2]
num_days_in_window = NUMBER_OF_DAYS_IN_WINDOW
points_per_day = NUM_DATA_POINTS_PER_DAY
points_in_window = num_days_in_window * points_per_day

# load testing data
testing = np.genfromtxt(test_file, delimiter=",")

# normalize
for row in testing:
    opening = row[0]
    for i in range(points_in_window):
```

```

row[i] = row[i] / opening

# pull X and y
X = np.array([p[range(points_in_window - 1)] for p in testing])
y = (testing[:, points_in_window - points_per_day + 1] > testing[:, points_in_window - points_per_day]) + 0

# test model
with open(model_file, "rb") as f:
    model = pickle.load(f)
    print model.score(X, y)

```

Python - sample.py (Randomly sampling training and testing data)

```

# Usage: python sample.py {NUM_TRAINING} {NUM_TESTING} {INPUT_FILE}

import pandas as pd
import numpy as np
import sys

if len(sys.argv) == 4:
    num_training = int(sys.argv[1])
    num_testing = int(sys.argv[2])
    input_file = sys.argv[3]

    training_file = "training%d_%s" % (num_training, input_file)
    testing_file = "testing%d_%s" % (num_testing, input_file)

    total_rows = num_training + num_testing
    data = pd.read_csv(input_file)
    if data.shape[0] < total_rows:
        print "Not enough rows in input file."

```

```

else:
    # randomly sample rows of input file
    rows = np.random.choice(data.index.values, total_rows)

    training_rows = sorted(rows[:num_training])
    testing_rows = sorted(rows[-num_testing:])

    # grab rows from dataframe
    data = np.asarray(data)
    training = data[training_rows]
    testing = data[testing_rows]

    np.savetxt(training_file, training, delimiter=",")
    np.savetxt(testing_file, testing, delimiter=",")
else:
    print "Incorrect usage. Use: \"python sample.py {NUM_TRAINING} {NUM_TESTING} {INPUT_FILE}\""

```

Visualization - index.html

```

<!DOCTYPE html>
<meta charset="utf-8">
<script src="http://d3js.org/d3.v3.js"></script>
<script src="http://code.jquery.com/jquery-1.8.2.js"></script>
<script src="http://code.jquery.com/ui/1.9.0/jquery-ui.js"></script>
<link href="http://code.jquery.com/ui/1.9.0/themes/cupertino/jquery-ui.css" rel="stylesheet" />
<link rel="stylesheet" type="text/css" href="/css/style.css">
<head>
  <script type="text/javascript" src="https://www.google.com/jsapi"></script>
  <script type="text/javascript">
    google.load("visualization", "1", {packages:["controls"]});
    google.load("visualization", "1", {packages:["table"]});
    google.load('visualization', '1', {packages:['annotationchart']});

```

```

var $symbols_path = "../data/sp500_symbols.txt"
var $dataset_path = "../data/2014-12-03_3yr_sp500_predictions/"
var $entries_per_stock = 7;

var sym_list = [];

$(document).ready(function() {
  // load all stock symbols from file
  var dsv = d3.dsv("\n", "text/plain");
  dsv($symbols_path, function(data){
    sym_list = data.map(function(d) {
      var sym = d["A"];

      // create a placeholder table div and chart div for each symbol
      $( ".container" ).append( $("<div class='sym_div' id='" + sym + "'>"
        + "<h4 class='symbol_title'>" + sym + "</h4>" +
        "<div class=chart_div id='" + sym + "'></div>" +
        "<div class=table_div id='" + sym + "'></div></div>"));

      return sym;
    })
  });

  google.setOnLoadCallback(drawAll);
});

// creates tables and charts for each stock
function drawAll() {
  // for each symbol, open it's csv file, and create a new table element
  sym_list.forEach(function(sym) {
    // read in the csv file
    d3.csv($dataset_path + sym + ".csv", function(data) {
      if(data != null && data != undefined) {
        // get relevant columns from dataset as an array of arrays
        var dataset = data.map(function(d) {

```

```

        if(d != undefined) {
            return [
                new Date(d["Date"]), +d["Open"],
                +d["Close"], Math.round(+d["Prediction"] * 1000) / 1000,
                $.parseJSON(d["Prediction Correct"].toLowerCase())
            ];
        } else {
            // remove the symbol tables if there is no data for them
            $('#'+ sym ).remove();

            return [];
        }
    });

    // add a table and chart for this symbol
    if(dataset != null) {
        drawTable(dataset, sym);
        drawChart(dataset, sym);
    } else {
        $('#'+ sym).remove();
    }
    } else {
        $('#'+ sym ).remove();
    }
    });
    });

    addSymbolFilter();
}

// set the search bar to filter divs by symbol
function addSymbolFilter() {
    var divs = $('#.sym_div');
    $('#search').on('keyup', function() {
        var val = $.trim(this.value);
        divs.hide();
        divs.filter(function() {

```

```

        return $(this).attr('id').search(val) >= 0
    }).show();
});
}

function drawTable(dataset, sym) {
    // initialize table with columns of dataset to use
    var dataTable = new google.visualization.DataTable();
    dataTable.addColumn('date', 'Date');
    dataTable.addColumn('number', 'Open');
    dataTable.addColumn('number', 'Close');
    dataTable.addColumn('number', 'Prediction');
    dataTable.addColumn('boolean', 'Prediction correct?');
    dataTable.allowHtml = true;

    // print the entries for this dataset
    var i;
    var pred_col = 4;

    for(i = 0; i < $entries_per_stock; ++i) {
        row = dataset[i];
        dataTable.addRow(row);

        if(row[pred_col] == true) {
            dataTable.setCell(i, pred_col, true, null,
                {'className': 'green-border center-text bold-green-font large-font'});
        }
    }

    // add the table to the div for this symbol
    var table = new google.visualization.Table(
        document.getElementById(sym).getElementsByClassName("table_div")[0]);
    table.draw(dataTable, {showRowNumber: true});
}

// add a chart for this symbol to show open vs close prices,

```

```

// with annotations that show prediction results
function drawChart(dataset, sym) {
  var dataTable = new google.visualization.DataTable();
  dataTable.addColumn('date', 'Date');
  dataTable.addColumn('number', 'Open');
  dataTable.addColumn('number', 'Close');
  dataTable.addColumn('string', 'Correct Prediction?');
  dataTable.addColumn('string', 'prediction value');
  // print the entries for this dataset
  var i;
  for(i = 0; i < $entries_per_stock; ++i) {
    row = dataset[i];
    // add a text entry for this row to use for annotation
    // description to show a string with correctness
    // (true/false) and the prediction itself
    row[4] = row[4].toString() + " : " + row[3];
    row.splice(3,1, "Correct Prediction?:");
    dataTable.addRow(row);
  }

  var options = {
    title: 'Opening vs Closing Prices',
    hAxis: {title: 'Date', titleTextStyle: {color: '#333'}},
    vAxis: {viewWindowMode:'pretty'},
    displayAnnotations: true
  };

  // add the chart to the div for this symbol
  var chart = new google.visualization.AnnotationChart(
    document.getElementById(sym).getElementsByClassName("chart_div")[0]);
  chart.draw(dataTable, options);
}
</script>
</head>

<body>

```

```
<label for="search" class="symbol_title" >Search By Symbol:</label>
<input type="search" name="filter" id="search" class="input-style-1"/>

<div class="container">
</div>
</body>
</html>
```

References:

1. <http://www.acit2k.org/ACIT/2013Proceedings/163.pdf>
2. <http://www.ccsenet.org/journal/index.php/mas/article/view/4586/3925>
3. <https://sites.google.com/site/predictingstockmovement/>
4. http://www.cs.berkeley.edu/~akar/EE671/report_stock.pdf